

Exercise 3: Normalizing Flow

Erik Stolt

What you did and how

We were tasked to implement three different types normalizing flow (NF); diagonal Gaussian, full 3D-Gaussian and full flow. Diagonal meaning that there is no off-diagonal elements in the covariance matrix when we assume Gaussian distributions of the variables. In full Gaussian we have off-diagonal elements and hence correlation. In full flow we have non-linear transformations g plus some affine (linear) transformation t in the end.

These NF:s were implemented using `Jammy Flow` which provides the skeleton for the NF. Then I implemented the CNN architecture to train the parameters of the NF to predict the parameters of the probability density function (pdf) of the labels T_{eff} , $\log g$ and $[Fe/h]$. In the end, we have a pdf of the labels (for a given input) which tells us "how certain" we are about the true value of the labels. Additionally, I presented the results in 5 main plots: Loss vs epoch, pdf visualization at different epochs, final distribution for the uncertainty, a joint pdf for two labels and histogram plots for each label also showing the true value for the input.

In order to obtain the target conditional distribution did I first create a grid and evaluated the grid points for the target pdf by calling `model.pdf(...)`. This then evaluated the target distribution on the grid, giving me the predicted pdf model.

To quantify my results, I plotted the coverage for the fraction of true values inside the model's 68th and 95th predicted intervals. For a properly trained model, the predicted percentiles should agree with the 68th and 95th percentile. This works even for non-Gaussian shapes, since the pdf should agree with the observations if trained properly. Moreover, this does not assume any specific shape for the pdf or the data.

What results you obtained

Overall, I obtained good results. Down below, I present some plots from the exercise. I trained all models for 100 with early stopping (patience = 10) and reduced learning rate on plateau (patience = 4). Additionally, I used batch size = 64. An example output during training looked like this:

```
Epoch 88/100 - Train Loss: 0.745930
- Val Loss: 0.938252 and lr: 3.91e-07
- No improvement for 3 epoch(s).
```

And some plots...




Figure 1: Loss for the training data, validation data and the test data over 100 epochs using diagonal Gaussian.



Figure 2: Loss for the training data, validation data and the test data over 100 epochs using full Gaussian.




Figure 3: Loss for the training data, validation data and the test data over 100 epochs using full flow.

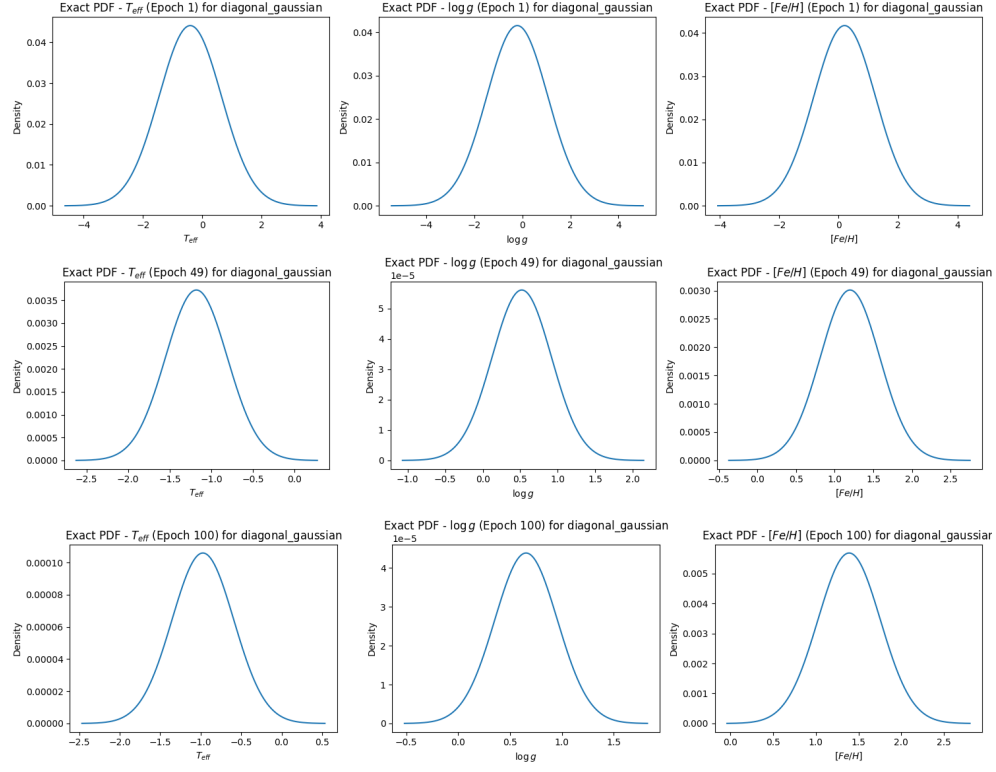


Figure 4: Evolution of the pdf for the three labels using diagonal Gaussian.

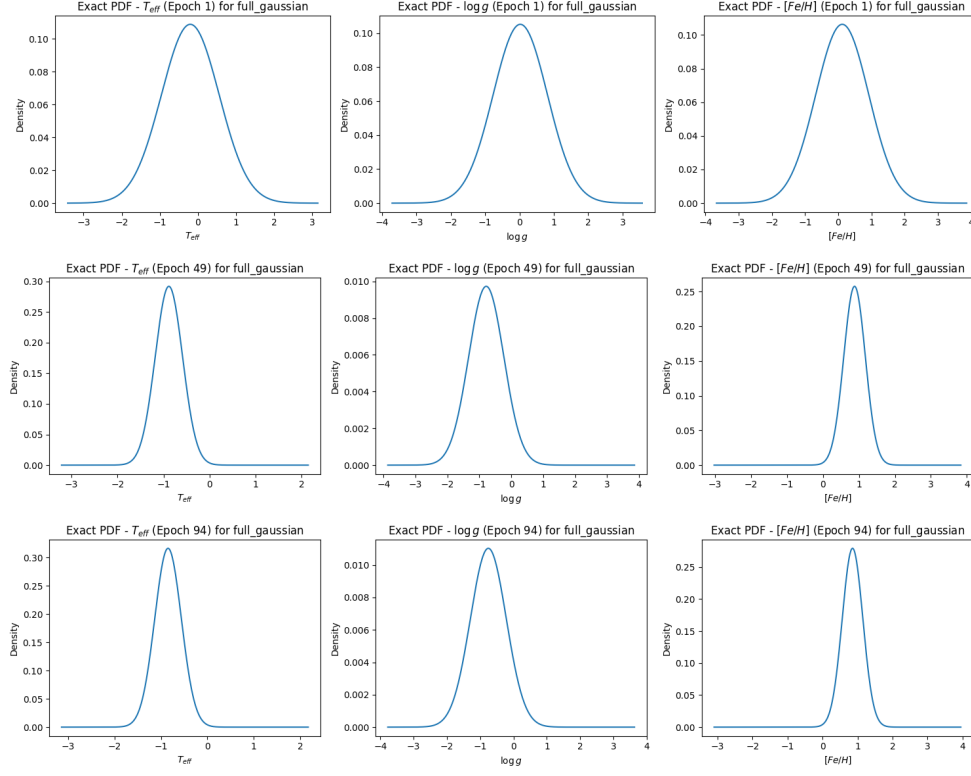


Figure 5: Evolution of the pdf for the three labels using full Gaussian.

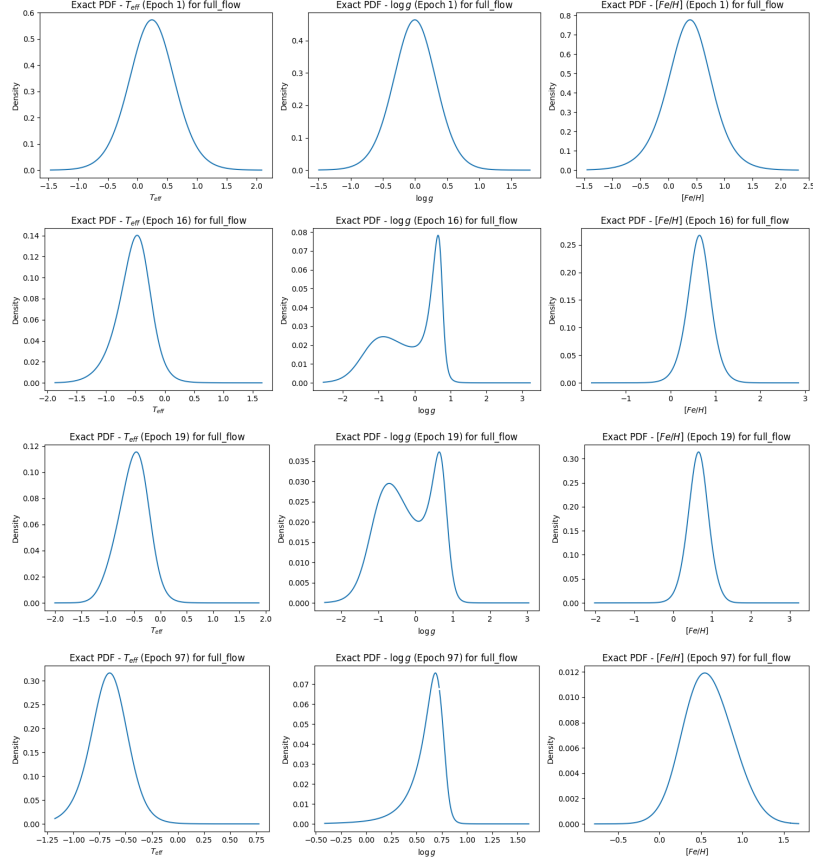


Figure 6: Evolution of the pdf for the three labels using full flow. We clearly see the non-linearity of the NF, especially in epochs 16 and 19. In the end the pdf's converge to somewhat of a Gaussian (though not quite).




Figure 7: Plot over the sampled target distribution for the three labels with a fitted Gaussian and the true value for the label. Here using diagonal Gaussian.




Figure 8: Plot over the sampled target distribution for the three labels with a fitted Gaussian and the true value for the label. Here using full Gaussian.

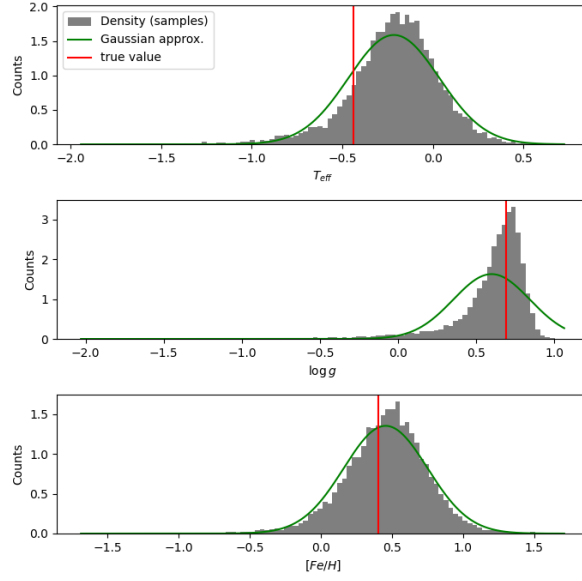


Figure 9: Plot over the sampled target distribution for the three labels with a fitted Gaussian and the true value for the label. Here using full flow.

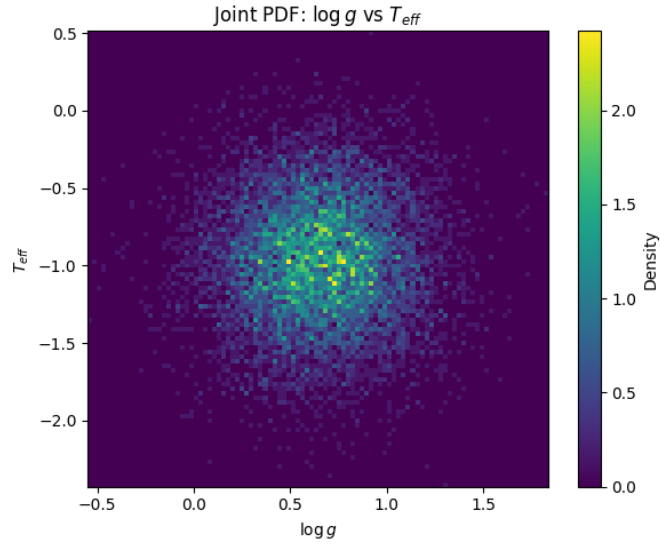


Figure 10: Heat map showing the joint pdf of $\log g$ and T_{eff} for diagonal Gaussian. Notice that there is no covariance.

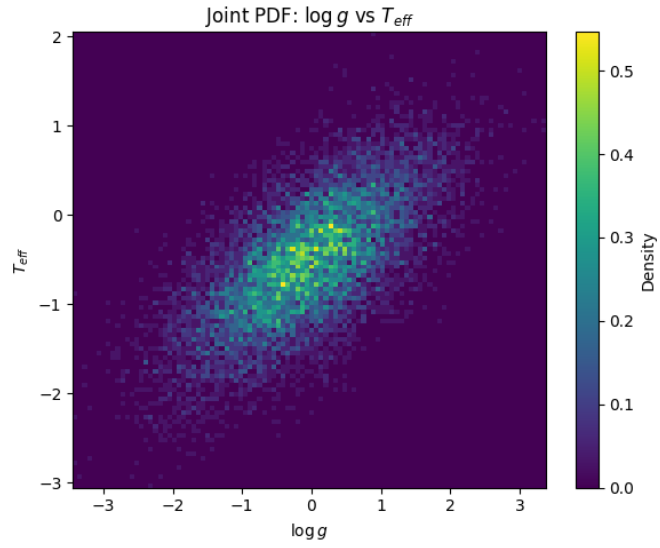


Figure 11: Heat map showing the joint pdf of $\log g$ and T_{eff} for full Gaussian. Notice that there is a covariance.




Figure 12: Heat map showing the joint pdf of $[Fe/H]$ and T_{eff} for full flow.

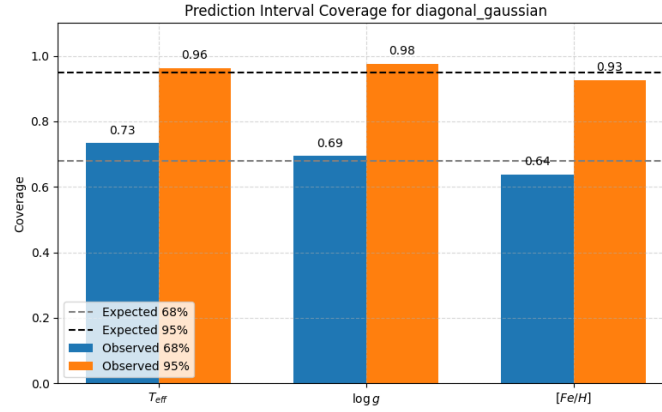


Figure 13: Convergence plot for each label for the 68th and 95th percentile for diagonal Gaussian.




Figure 14: Convergence plot for each label for the 68th and 95th percentile for full Gaussian.

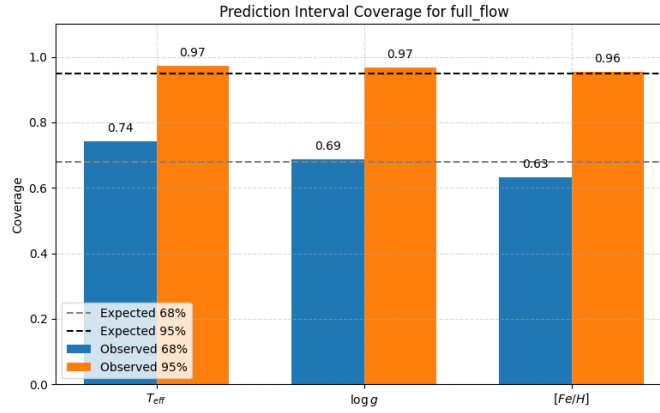


Figure 15: Convergence plot for each label for the 68th and 95th percentile for full flow.

Remark: The plots were plotted using the first element in the first batch in the test set with random seed = 42 (see more in discussion), and therefore not fully representative for the entire dataset. This however was not the exercise, only to plot some pdf's and to predict the pdf for a given input spectra.

Discussion

The first thing i noticed when training the network was that dropout really did not work. This indicates to me that the model is very sensitive to perturbing/dropping nodes which results in poor parameter predictions. Therefore, I trained the model without dropout.

We observe that the model's predicted uncertainties are not always centered on the true values. This is, however, not surprising when looking at only a single batch. The model predicts a full

probability distribution conditioned on the input, and the mean of this distribution may differ from the true label for individual examples. Importantly, we do not expect each individual prediction to be unbiased. However, if the model is well-calibrated and unbiased overall, the average residual (i.e., prediction minus true value) across the full test set should be approximately zero. We strengthen this by looking in the convergence plot. There we see that percentiles, sampled from the conditional distribution, matches well with the observed percentiles. When a data point is far from the mean, the model's uncertainty (assuming it's well-trained) tends to predict a larger correction for that point. This means that the pdf may not be centered around the true mean of the distribution, as the model is adjusting the data point towards the mean. Essentially, the larger corrections for outlier points cause the pdf to shift, reflecting the model's attempts to correct those points.

The reason for full flow converging to a almost Gaussian shape means that the underlying data already is distributed similar to a Gaussian. We see that the model is indeed capable of predicting more complex patterns but since it converges and stays roughly the same for 30 epochs means that its confident in its shape. Which is further strengthen by the convergence plot.

The coverage plot for full Gaussian shows that the model is well-calibrated at the 95% level across all parameters but slightly underestimates uncertainty at the 68% level, particularly for T_{eff} and $[\text{Fe}/\text{H}]$. This behavior is expected for a *full Gaussian* model, which may struggle to capture asymmetries or heavier tails in the true posterior distribution. A more expressive model, such as a full flow, could better capture the distributional shape and improve coverage at narrower confidence levels, which we indeed see!

I also chose to keep the fitted gaussian in the visualization plots even for full flow since it was given in the template and one can see it does not fit very well compared to the full flow NF. Lastly, we see that full flow captures more complex patterns and that the diagonal Gaussian only has a simple 3D Gaussian with no correlation between the labels. Full Gaussian is also a 3D Gaussian but with correlation.

What challenges you encountered and what could be improved

The main challenge was that the data uncertainties matched very well with a simple Gaussian, hence, the full flow did not deviate a lot from the Gaussian. This meant that I had to train the network a lot in order to see a more complex pdf. However, this was resolved in the later stage of the exercise where I saw that I had incorrectly split the data. When this was resolved much better results and nicer pdf's were obtained!

Finding a good network architecture is always crucial in ML problems. I have tried a few networks, but there is probably a better one. If I had more time I would try some more networks. Moreover, It was also somewhat difficult to fully understand how **Jammy Flows** really works. This led me to dedicate a lot of time just to understand how the template code worked. In the end, I would say I have a good grasp of how it works!

Link to Github

Link: <https://github.com/erst6955/Advanced-Applied-Deep-Learning-in-Physics-And-Engineering>

Code (for backup)

Remark: The only thing I changed between the final runs was the pdf model. Everything else was unchanged.

~\Desktop\Advanced Applied Deep Learning In Physics And Engineering\Exercises\Exercise 3 -
Normalizing Flows\stellar_prediction.py

```
1 import time
2 import sys
3 import os
4 import argparse
5 import glob
6 import subprocess
7 import numpy as np
8 import torch
9 import torch.nn as nn
10 import torch.optim as optim
11 from torch.utils.data import DataLoader, TensorDataset, random_split
12 from matplotlib import pyplot as plt
13 import jammy_flows
14 from scipy.stats import norm
15 from sklearn.preprocessing import MinMaxScaler
16 from torch.utils.data import Dataset, DataLoader, random_split
17 import pylab
18 from sklearn.preprocessing import StandardScaler
19 import random
20
21
22 # Set random seed for reproducibility
23 seed = 42
24 torch.manual_seed(seed)
25 random.seed(seed)
26 np.random.seed(seed)
27 torch.manual_seed(seed)
28 if torch.cuda.is_available():
29     torch.cuda.manual_seed_all(seed)
30
31
32 DATA_PATH = r"C:\Users\Erik\Desktop\Advanced Applied Deep Learning In Physics And
Engineering\Datasets\Astronomy"
33 fp64_on_cpu = False
34
35 # Hyperparameters
36 #learning_rate = 1e-4
37
38 names = ["$T_{\text{eff}}$", "$\log g$", "$[Fe/H]$"]
39
40
41
42 # ===== Jammy Flow functions =====
43
44 # Defining the normalizing flow model is a bit more involved and requires knowledge of the
jammy_flows library.
45 # Therefore, we provide the relevant code here.
```

```

46 class CombinedModel(nn.Module):
47     """
48     A combined model that integrates a normalizing flow with a CNN encoder.
49     """
50
51     def __init__(self, encoder, nf_type="diagonal_gaussian"):
52         """
53         Initializes the normalizing flow model.
54
55         Parameters
56         -----
57         encoder : callable
58             A function or callable object that returns an encoder model. The encoder model
59             should take the number of flow parameters as input and output the latent
60             dimension.
61         nf_type : str, optional
62             The type of normalizing flow to use. Options are "diagonal_gaussian",
63             "full_gaussian",
64             and "full_flow". Default is "diagonal_gaussian".
65         Raises
66         -----
67         Exception
68             If an unknown `nf_type` is provided.
69         Notes
70         -----
71             This method sets up a 3-dimensional probability density function (PDF) over Euclidean
72             space (e3)
73             using the specified normalizing flow type. The flow structure and options are
74             configured based on
75             the provided `nf_type`. The PDF is created using the `jammy_flows` library, and the
76             number of flow
77             parameters is determined and printed. The encoder is initialized with the number of
78             flow parameters.
79         """
80
81         super().__init__()
82
83         # we define a 3-d PDF over Euclidean spae (e3)
84         # using recommended settings (https://github.com/thoglu/jammy\_flows/issues/5 scroll
85         down)
86
87         opt_dict = {}
88         opt_dict["t"] = {}
89         if (nf_type == "diagonal_gaussian"):
90             opt_dict["t"]["cov_type"] = "diagonal"
91             flow_defs = "t"
92         elif (nf_type == "full_gaussian"):
93             opt_dict["t"]["cov_type"] = "full"
94             flow_defs = "t"
95         elif (nf_type == "full_flow"):
96             opt_dict["t"]["cov_type"] = "full"
97             flow_defs = "gggt"

```



```

90         else:
91             raise Exception("Unknown nf type ", nf_type)
92
93         opt_dict["g"] = dict()
94         opt_dict["g"]["fit_normalization"] = 1
95         opt_dict["g"]["upper_bound_for_widths"] = 1.0
96         opt_dict["g"]["lower_bound_for_widths"] = 0.01
97
98         self.nf_type = nf_type
99
100         # 3d PDF (e3) with ggggt flow structure. Four Gaussianation-flow
101         # (https://arxiv.org/abs/2003.01941) layers ("g") and an affine flow ("t")
102         self.pdf = jammy_flows.pdf("e3", flow_defs, options_overwrite=opt_dict,
103                                   amortize_everything=True, amortization_mlp_use-
104                                   _custom_mode=True)
105
106         # get the number of flow parameters based on the selected model
107         num_flow_parameters = self.pdf.total_number_amortizable_params
108
109         print("The normalizing flow has ", num_flow_parameters, " parameters...")
110
111         # latent dimension (output of the CNN encoder) is set to 128
112         self.encoder = encoder(num_flow_parameters)
113
114     def log_pdf_evaluation(self, target_labels, input_data):
115         """
116         Evaluate the log probability density function (PDF) for the given target labels and
117         input data.
118
119         The normalizing flow parameters are predicted by the encoder network based on the
120         input data.
121         Then, the log PDF is evaluated at the position of the label.
122
123         Parameters:
124         -----
125         target_labels : torch.Tensor
126             The target labels for which the log PDF is to be evaluated.
127         input_data : torch.Tensor
128             The input data to be encoded and used for evaluating the log PDF.
129         Returns:
130         -----
131         log_pdf : torch.Tensor
132             The evaluated log PDF for the given target labels and input data.
133         """
134         latent_intermediate = self.encoder(input_data) # get trained flow parameters from
135         # the CNN encoder
136
137         if (self.nf_type == "full_flow"):
138             # convert to double. Double precision is needed for the Gaussianization flow.
139             # This is for numerical stability.

```

```

134         if fp64_on_cpu: # MPS does not support double precision, therefore we need to
run the flow on the CPU
135             latent_intermediate = latent_intermediate.cpu().to(torch.float64)
136             target_labels = target_labels.cpu().to(torch.float64)
137         else:
138             latent_intermediate = latent_intermediate.to(torch.float64)
139             target_labels = target_labels.to(torch.float64)
140
141         # evaluate the log PDF at the target labels. We use log pdf for numerical stability.
142         log_pdf, _, _ = self.pdf(target_labels, amortization_parameters=latent_intermediate)
143         return log_pdf
144
145     def sample(self, flow_params, samplesize_per_batchitem=1000):
146         """
147         Sample new points from the PDF given input data.
148
149         Parameters
150         -----
151         flow_params : tensor
152             Parameters for the normalizing flow, must be of shape (B, L) where B is the batch
size and L is the latent dimension.
153         samplesize_per_batchitem : int, optional
154             Number of samples to draw per batch item. Defaults to 1000.
155
156         Returns
157         -----
158         tensor
159             A tensor of shape (B, S, D) where B is the batch dimension, S is the number of
samples,
160             and D is the dimension of the target space for the samples.
161         """
162         # for full flow we need to convert to double precision for the normalizing flow
163         # for numerical stability
164         if (self.nf_type == "full_flow"):
165             # convert to double
166             if fp64_on_cpu: # MPS does not support double precision, therefore we need to run
the flow on the CPU
167                 flow_params = flow_params.cpu().to(torch.float64)
168             else:
169                 flow_params = flow_params.to(torch.float64)
170
171             batch_size = flow_params.shape[0] # get the batch size
172             # sample from the normalizing flow
173             repeated_samples, _, _, _ = self.pdf.sample(amortization_paramet-
ers=flow_params.repeat_interleave(
174                 samplesize_per_batchitem, dim=0), allow_gradients=False)
175
176             # reshape the samples to be grouped by batch item
177             reshaped_samples = repeated_samples[:, None, :].view(
178                 batch_size, samplesize_per_batchitem, -1)
179

```

```

180         return reshaped_samples
181
182     def forward(self, input_data, samplesize_per_batchitem=1000):
183         """
184         Perform a forward pass through the model, predicting the mean and standard deviation
185         of the samples.
186
187         Normalizing flows do not directly predict the target labels. Instead, they predict
188         the parameters of the flow that
189         transforms the base distribution to the target distribution. Often, we still want to
190         predict the target labels.
191         Then, we can sample from the distribution and form the mean of the samples and their
192         standard deviations.
193         This is what this function does.
194
195         Parameters
196         -----
197         input_data : torch.Tensor
198             The input data tensor.
199         Returns
200         -----
201         torch.Tensor
202             A tensor of size (B, D*2) where the first half (size D) are the means,
203             the second half (another D) are the standard deviations.
204         """
205         flow_params=self.encoder(input_data)
206         samples=self.sample(flow_params, samplesize_per_batchitem=samplesize_per_batchitem)
207
208         # form mean along dim 1 (samples)
209         means=samples.mean(dim=1)
210         # form std along dim 1 (samples)
211         std_deviations=samples.std(dim=1)
212
213         # return means and std deviations as a concatenated tensor along dim 1
214         return torch.cat([means, std_deviations], dim=1)
215
216     def visualize_pdf(self, input_data, filename, pdf_model, samplesize=10000, batch_index=0,
217                       truth=None):
218         """
219         Visualizes the probability density function (PDF) of the given input data using a
220         normalizing flow model.
221
222         The function generates samples from the normalizing flow (using the sample()
223         function)
224         and plots the histogram of the samples together with a Gaussian approximation.
225
226         Parameters
227         -----
228         input_data : torch.Tensor
229             The input data tensor from which to pick one batch item for visualization.
230         filename : str

```

```

224         The filename where the resulting plot will be saved.
225     samplesize : int, optional
226         The number of samples to generate for the PDF visualization (default is 10000).
227     batch_index : int, optional
228         The index of the batch item to visualize (default is 0).
229     truth : torch.Tensor, optional
230         The true values of the labels, used for comparison in the plot (default is None).
231
232     Returns
233     -----
234     None
235     """
236     # pick out one input from batch
237     input_bitem = input_data[batch_index:batch_index+1]
238
239     # get the flow parameters (by passing the input data through the CNN encoder
network). This is basically evaluation of the encoder network.
240     flow_params = self.encoder(input_bitem)
241
242     # sample from the normalizing flow (i.e. samples are drawn from the base distribution
and transformed by the flow
243     # using the change-of-variable formula)
244     samples = self.sample(flow_params, samplesize_per_batchitem=samplesize)
245     # the rest of the code is just plotting.
246
247     # we only have 1 batch item
248     samples = samples.squeeze(0)
249
250     # plot three 1-dimensional distributions together with normal approximation,
251     # so we calculate the mean and std of the samples
252     mean = samples.mean(dim=0).cpu().numpy()
253     std = samples.std(dim=0).cpu().numpy()
254     samples = samples.cpu().numpy()
255
256     fig, axdict = plt.subplots(3, 1, figsize=(8, 8))
257
258     fig.subplots_adjust(hspace=0.4) # Adjust the space between subplots
259     for dim_ind in range(3):
260         ax = axdict[dim_ind]
261         ax.hist(samples[:, dim_ind], color="k", density=True, bins=100, alpha=0.5,
label="Density (samples)")
262
263         # Gaussian overlay
264         xvals = np.linspace(samples[:, dim_ind].min(), samples[:, dim_ind].max(), 1000)
265         yvals = norm.pdf(xvals, loc=mean[dim_ind], scale=std[dim_ind])
266         ax.plot(xvals, yvals, color="green", label="Gaussian approx.")
267
268         # plot the true value if it is given
269         if (truth is not None):
270             true_value = truth[dim_ind].cpu().item()
271             axdict[dim_ind].axvline(

```

```

272         true_value, color="red", label="true value")
273
274     # Add axis labels
275     axdict[dim_ind].set_xlabel(names[dim_ind]) # x-axis label
276     axdict[dim_ind].set_ylabel("Counts") # y-axis label
277
278     # plot the legend only for the first panel
279     if (dim_ind == 0):
280         axdict[dim_ind].legend()
281
282     plt.savefig(f"{filename}_{pdf_model}") # Save the plot to a file
283     plt.show() # Display the plot
284     plt.close(fig)
285
286
287     # === Add 2D heatmap plot for joint distribution ===
288     dim1, dim2 = 1, 0 # Example: Teff vs log g
289     fig2, ax2 = plt.subplots(figsize=(6, 5))
290     h = ax2.hist2d(samples[:, dim1], samples[:, dim2], bins=100, cmap='viridis',
291 density=True, edgecolors='none')
292     plt.colorbar(h[3], ax=ax2, label="Density")
293
294     ax2.set_xlabel(names[dim1])
295     ax2.set_ylabel(names[dim2])
296     ax2.set_title(f"Joint PDF: {names[dim1]} vs {names[dim2]}")
297
298     # if truth is not None:
299     #     ax2.plot(truth[dim1].cpu().item(), truth[dim2].cpu().item(), "rx", label="True
300 value")
301     #     ax2.legend()
302
303     plt.tight_layout()
304     fig2.savefig(f"joint_pdf_heatmap_{pdf_model}")
305     plt.show()
306     plt.close(fig2)
307
308
309 # ===== Data related functions =====
310 def get_normalized_data(data_path):
311     spectra = np.load(f"{data_path}\\spectra.npy")
312     spectra_length = spectra.shape[1]
313     # labels: mass, age, l_bol, dist, t_eff, log_g, fe_h, SNR
314     labelNames = ["mass", "age", "l_bol", "dist", "t_eff", "log_g", "fe_h", "SNR"]
315     labels = np.load(f"{data_path}\\labels.npy")
316
317     # We only use the three labels: t_eff, log_g, fe_h
318     labelNames = labelNames[-4:-1]
319     labels = labels[:, -4:-1]

```

```

320     n_labels = labels.shape[1]
321
322     # normalize the spectra and labels via log
323     spectra = np.log(np.maximum(spectra, 0.2))
324
325     # scale all labels with minmaxscaler independently and keep the parameters for unscaling.
326
327     scaler = StandardScaler()
328     labels = scaler.fit_transform(labels)
329
330     print("Spectra shape: ", spectra.shape)
331
332     return spectra, labels, spectra_length, n_labels, labelNames, scaler
333
334
335 def get_datasets(spectra, labels, split_ratio=0.1, batch_size=64):
336     """
337     Create datasets for training, validation, and testing.
338
339     Returns
340     -----
341     tuple
342         A tuple containing the training, validation, and test datasets.
343     """
344     # Create a TensorDataset from the data
345     spectra_tensor = spectra.clone().detach() # Use clone().detach() to avoid warnings.
346     labels_tensor = labels.clone().detach()
347     dataset = TensorDataset(spectra_tensor, labels_tensor)
348
349     n = len(dataset)
350     val_size = int(split_ratio * n)
351     test_size = int(split_ratio * n)
352     train_size = n - val_size - test_size # Ensure full coverage
353
354     train_dataset, val_dataset, test_dataset = random_split(dataset, [train_size, val_size,
355     test_size])
356
357     print(f"Train size: {len(train_dataset)}, Validation size: {len(val_dataset)}, Test size:
358     {len(test_dataset)}")
359
360     train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True) # create
361     val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
362     test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
363
364     return train_loader, val_loader, test_loader
365
366 # ===== Model related functions =====
367 def plot_loss(train_losses, val_losses, test_loss, pdf_model, filename="loss_plot.png"):

```

```

368     """
369     Plots the training, validation, and test loss.
370
371     Parameters
372     -----
373     train_losses : list
374         List of training loss values for each epoch.
375     val_losses : list
376         List of validation loss values for each epoch.
377     test_loss : float
378         The test loss value.
379     filename : str, optional
380         The filename where the plot will be saved (default is "loss_plot.png").
381     """
382     plt.figure(figsize=(10, 6))
383     plt.plot(range(1, len(train_losses) + 1), train_losses, label="Training Loss",
384             marker="o")
385     plt.plot(range(1, len(val_losses) + 1), val_losses, label="Validation Loss", marker="o")
386     plt.axhline(y=test_loss, color="red", linestyle="--", label="Test Loss")
387     plt.xlabel("Epoch")
388     plt.ylabel("Loss")
389     plt.title("Training, Validation, and Test Loss")
390     plt.legend()
391     plt.grid(True, linestyle="--", linewidth=0.5)
392     plt.savefig(f"{filename}_{pdf_model}")
393     plt.show()
394     plt.close()
395
396 def train_nf_model(model, train_loader, val_loader, device, pdf_model, num_epochs=50,
397                   learning_rate=1e-4, num_grid=1000):
398     """
399     Trains the normalizing flow model.
400
401     Parameters
402     -----
403     model : nn.Module
404         The combined normalizing flow and encoder model.
405     train_loader : DataLoader
406         DataLoader for the training data.
407     val_loader : DataLoader
408         DataLoader for the validation data.
409     device : torch.device
410         The device to run training on (CPU, CUDA, or MPS).
411     num_epochs : int
412         Number of epochs to train.
413     learning_rate : float
414         Learning rate for the optimizer.
415
416     Returns

```

```

416     -----
417     model : nn.Module
418         The trained model.
419     train_losses : list
420         List of training loss values for each epoch.
421     val_losses : list
422         List of validation loss values for each epoch.
423     """
424     best_val_loss = float("inf")
425     epochs_without_improvement = 0
426     early_stopping_patience = 10
427     reduce_lr_patience = 4
428
429     optimizer = optim.Adam(model.parameters(), lr=learning_rate)
430     scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5,
patience=reduce_lr_patience, min_lr=1e-8, verbose=True)
431
432     train_losses = []
433     val_losses = []
434
435
436     for epoch in range(num_epochs):
437         model.train()
438         running_loss = 0.0
439
440         for batch_inputs, batch_labels in train_loader:
441             batch_inputs, batch_labels = batch_inputs.to(device), batch_labels.to(device) #
move data to device
442
443             optimizer.zero_grad() # zero the gradients
444             loss = nf_loss(batch_inputs, batch_labels, model) # compute the loss
445             loss.backward() # update the gradients
446             optimizer.step() # update the weights
447
448             running_loss += loss.item() * batch_inputs.size(0) # accumulate the loss wighted
by the batch size
449
450             avg_train_loss = running_loss / len(train_loader.dataset) # average the loss over the
dataset
451             train_losses.append(avg_train_loss)
452             print(f"Epoch {epoch+1}/{num_epochs} - Train Loss: {avg_train_loss:.6f}")
453
454             # Optional: Validation loop
455             if val_loader is not None:
456                 model.eval()
457                 val_loss = 0.0
458                 with torch.no_grad():
459                     for val_inputs, val_labels in val_loader:
460                         val_inputs, val_labels = val_inputs.to(device), val_labels.to(device)
461                         val_loss += nf_loss(val_inputs, val_labels, model).item() *
val_inputs.size(0)

```



```

462
463         avg_val_loss = val_loss / len(val_loader.dataset)
464         val_losses.append(avg_val_loss)
465
466         for param_group in optimizer.param_groups:
467             print(f"                - Val Loss:    {avg_val_loss:.6f} and lr:
468 {param_group['lr']:.2e}")
469
470         scheduler.step(avg_val_loss) # Reduce learning rate on plateau
471
472
473
474         # ===== Plot the PDF for a test example every 3 epochs ===== Here we create a
475         automatic range for the grid and evaluate the target pdf
476         if epoch % 3 == 0: # # Plot the PDF for a test example every 3 epochs
477             model.eval() # set the model to evaluation mode
478             with torch.no_grad(): # no gradients needed
479                 # Grab one test example
480                 test_input, test_label = next(iter(test_loader)) # get the first batch of the
481                 test set
482                 x_example = test_input[0].unsqueeze(0).to(device) # add batch dimension (s.t.
483                 [B, 1, 16 384]) and move to device
484                 truth = test_label[0].cpu().numpy() # get the true labels for the first batch
485                 item
486
487                 fig, axes = plt.subplots(1, 3, figsize=(15, 4)) # create a figure with 3
488                 subplots, one for each label
489                 for i, label_name in enumerate(["T_{eff}$", "$\log g$", "$[Fe/H]$"]):
490                     # Sample from the flow to estimate mean and std
491                     samples = model.sample(model.encoder(x_example), samplesize_per_batch-
492                     item=1000)[0].cpu().numpy() # pass the input through the encoder (i.e. the CNN) to get the
493                     flow parameters
494
495                     # and then sample from the target distribution.
496                     mean = samples[:, i].mean() # get the mean of the samples for each label
497                     std = samples[:, i].std() # get the std of the samples for each label
498
499                     # Create grid: mean  $\pm$  6 stds
500                     xvals = np.linspace(mean - 4 * std, mean + 4 * std, num_grid) # create a
501                     grid of x values for the i-th label
502                     grid = np.zeros((num_grid, 3)) # create a grid of 500 points for each
503                     label with mean at 0.5 (doesn't really matter, we will overwrite it)
504                     grid[:, i] = xvals # set the i-th label to the xvals
505
506                     grid_tensor = torch.tensor(grid, dtype=torch.float32,
507                     device=x_example.device) # create a tensor from the grid and move to device
508                     flow_params = model.encoder(x_example) # get the flow parameters from the
509                     encoder
510
511                     log_pdf, _, _ = model.pdf(grid_tensor, amortization_paramet-
512                     ers=flow_params.expand(num_grid, -1)) # evaluate the target distribution log pdf at the grid
513                     points

```

```

499         pdf = torch.exp(log_pdf).cpu().numpy() # get the pdf by exponentiating
the log pdf
500
501         # Plotting
502         axes[i].plot(xvals, pdf)
503         #axes[i].axvline(truth[i], color='red', linestyle='--', label='True')
504         axes[i].set_xlabel(label_name)
505         axes[i].set_ylabel("Density")
506         axes[i].set_title(f"Exact PDF - {label_name} (Epoch {epoch+1}) for
{pdf_model}")
507         #axes[i].legend()
508
509         plt.tight_layout()
510         plt.savefig(f"exact_pdf_epoch_{epoch+1}_{pdf_model}.png")
511         plt.close()
512
513
514
515
516         # ===== Early stopping based on validation loss =====
517         if avg_val_loss < best_val_loss - 1e-5:
518             best_val_loss = avg_val_loss
519             epochs_without_improvement = 0
520         else:
521             epochs_without_improvement += 1
522             print(f"                - No improvement for {epochs_without_improvement}
epoch(s).")
523
524             if epochs_without_improvement >= early_stopping_patience:
525                 print("Early stopping triggered.")
526                 break
527
528         return model, train_losses, val_losses
529
530
531
532     def evaluate_test_loss(model, test_loader, device):
533         """
534         Evaluates the model on the test dataset and computes the test loss.
535
536         Parameters
537         -----
538         model : nn.Module
539             The trained model.
540         test_loader : DataLoader
541             DataLoader for the test data.
542         device : torch.device
543             The device to run evaluation on (CPU, CUDA, or MPS).
544
545         Returns
546         -----

```

```

547     float
548         The computed test loss.
549     """
550     model.eval()
551     test_loss = 0.0
552     with torch.no_grad():
553         for test_inputs, test_labels in test_loader:
554             test_inputs, test_labels = test_inputs.to(device), test_labels.to(device)
555             test_loss += nf_loss(test_inputs, test_labels, model).item() *
test_inputs.size(0)
556     test_loss /= len(test_loader.dataset)
557
558     return test_loss
559
560
561
562 # Define the CNN encoder model. The output of the model is the input to the normalizing flow.
563 # The latent dimension is the number of parameters in the normalizing flow.
564 class TinyCNNEncoder(nn.Module):
565     def __init__(self, latent_dimension):
566         super().__init__()
567         self.encoder = nn.Sequential(
568             nn.Conv1d(1, 8, kernel_size=7, stride=1, padding=3),      # [B, 8, 16384]
569             nn.BatchNorm1d(8),
570             nn.ReLU(),
571             nn.MaxPool1d(kernel_size=4),                            # [B, 8, 4096]
572
573             nn.Conv1d(8, 16, kernel_size=5, stride=1, padding=2),   # [B, 16, 4096]
574             nn.BatchNorm1d(16),
575             nn.ReLU(),
576             nn.MaxPool1d(kernel_size=4),                            # [B, 16, 1024]
577
578             nn.Conv1d(16, 32, kernel_size=5, stride=1, padding=2),  # [B, 32, 1024]
579             nn.BatchNorm1d(32),
580             nn.ReLU(),
581             nn.MaxPool1d(kernel_size=4),                            # [B, 32, 256]
582
583             nn.Conv1d(32, 64, kernel_size=3, stride=1, padding=1),  # [B, 64, 256]
584             nn.BatchNorm1d(64),
585             nn.ReLU(),
586             nn.AdaptiveAvgPool1d(1),                                # [B, 64, 1]
587         )
588
589         self.project = nn.Sequential(
590             nn.Flatten(),                                           # [B, 64]
591             nn.Linear(64, 256),
592             nn.ReLU(),
593             nn.Linear(256, latent_dimension)
594         )
595

```

```

596
597
598     def forward(self, x):
599         x = self.encoder(x)          # Expect input shape: [B, 1, 16384]
600         x = self.project(x)          # Output shape: [B, latent_dim]
601         return x
602
603
604 def nf_loss(inputs, batch_labels, model):
605     """
606     Computes the loss for a normalizing flow model, according to maximum likelihood
607     estimation.
608     The loss is defined as the negative log probability of the labels given the input data.
609     This loss
610     allows the model to learn the true value of the normalizing flow parameters.
611
612     Parameters
613     -----
614     inputs : torch.Tensor
615         The input data to the model.
616     batch_labels : torch.Tensor
617         The labels corresponding to the input data.
618     model : torch.nn.Module
619         The normalizing flow model used for evaluation.
620     Returns
621     -----
622     torch.Tensor
623         The computed loss value.
624     """
625     log_pdfs = model.log_pdf_evaluation(batch_labels, inputs) # get the probability of the
626     labels given the input data
627     loss = -log_pdfs.mean() # take the negative mean of the log probabilities
628     return loss
629
630 # ===== Quantifying Model Functions =====
631 def compute_and_plot_coverage(model, test_loader, device, scaler, pdf_model, n_samples=1000):
632     """
633     Computes and plots the coverage of the predicted intervals for the test dataset.
634     """
635
636     model.eval()
637     all_true = []
638     all_samples = []
639
640     with torch.no_grad():
641         for batch_inputs, batch_labels in test_loader: # get the test data
642             batch_inputs = batch_inputs.to(device) # move to device
643             batch_labels = batch_labels.to(device)

```

```

644
645         flow_params = model.encoder(batch_inputs) # get the flow parameters from the
encoder (i.e. the CNN)
646         samples = model.sample(flow_params, samplesize_per_batchitem=n_samples) # sample
from the normalizing flow from each pdf using the flow parameters
647
648         all_true.append(batch_labels.cpu().numpy()) # [B, n_labels]
649         all_samples.append(samples.cpu().numpy()) # [B, n_samples, 3]
650
651     y_true = np.concatenate(all_true, axis=0) # true labels
652     y_samples = np.concatenate(all_samples, axis=0) # sampled labels
653
654     # === Denormalize using StandardScaler ===
655     y_true = scaler.inverse_transform(y_true)
656     y_samples = y_samples * scaler.scale_[None, None, :] + scaler.mean_[None, None, :]
657
658     coverage_68 = []
659     coverage_95 = []
660
661     # Calculate coverage for each label
662     for i in range(y_true.shape[1]):
663         lower_68 = np.percentile(y_samples[:, :, i], 16, axis=1)
664         upper_68 = np.percentile(y_samples[:, :, i], 84, axis=1)
665         coverage_68_i = ((y_true[:, i] >= lower_68) & (y_true[:, i] <= upper_68)).mean() #
Fraction of true values inside the model's 68% predicted interval
666         coverage_68.append(coverage_68_i)
667
668         lower_95 = np.percentile(y_samples[:, :, i], 2.5, axis=1)
669         upper_95 = np.percentile(y_samples[:, :, i], 97.5, axis=1)
670         coverage_95_i = ((y_true[:, i] >= lower_95) & (y_true[:, i] <= upper_95)).mean() #
Fraction of true values inside the model's 95% predicted interval
671         coverage_95.append(coverage_95_i)
672
673     # === Plot ===
674     x = np.arange(len(names))
675     width = 0.35
676
677     fig, ax = plt.subplots(figsize=(8, 5))
678     bars1 = ax.bar(x - width/2, coverage_68, width, label="Observed 68%")
679     bars2 = ax.bar(x + width/2, coverage_95, width, label="Observed 95%")
680
681     ax.axhline(0.68, color='gray', linestyle='--', label="Expected 68%")
682     ax.axhline(0.95, color='black', linestyle='--', label="Expected 95%")
683
684
685     # Add values on top of bars
686     for bar in bars1:
687         height = bar.get_height()
688         ax.text(bar.get_x() + bar.get_width()/2., height + 0.02,
689               f'{height:.2f}', ha='center', va='bottom', fontsize=10)
690

```

```

691     for bar in bars2:
692         height = bar.get_height()
693         ax.text(bar.get_x() + bar.get_width()/2., height + 0.02,
694               f'{height:.2f}', ha='center', va='bottom', fontsize=10)
695
696
697     ax.set_ylabel("Coverage")
698     ax.set_ylim(0, 1.1)
699     ax.set_xticks(x)
700     ax.set_xticklabels(names)
701     ax.set_title(f"Prediction Interval Coverage for {pdf_model}")
702     ax.legend()
703     ax.grid(True, linestyle="--", alpha=0.5)
704
705     plt.tight_layout()
706     plt.savefig(f"coverage_plot_percentiles_{pdf_model}")
707     plt.show()
708
709
710
711 # ===== Run Code =====
712
713 # Parameters
714 epochs = 100
715 pdf_model = "diagonal_gaussian" # choose complexity of the resulting pdf --> more parameters
716
717
718 if __name__ == "__main__":
719
720     # ===== Setup =====
721     parser = argparse.ArgumentParser() ## Create an argument parser
722     parser.add_argument("-normalizing_flow_type", default=pdf_model,
723                       choices=["diagonal_gaussian", "full_gaussian", "full_flow"]) # Add an
724     argument for the normalizing flow type
725     args = parser.parse_args() # Parse the arguments
726     print("Using normalizing flow type ", args.normalizing_flow_type)
727
728     model = CombinedModel(TinyCNNEncoder, nf_type=args.normalizing_flow_type) # Create the
729     model with the specified normalizing flow type
730
731     # Detect and use Apple Silicon GPU (MPS) if available
732     device = torch.device("cuda" if torch.cuda.is_available() else "mps" if
733 torch.backends.mps.is_available() else "cpu")
734     if args.normalizing_flow_type == "full_flow" and device.type == "mps":
735         # MPS does not support double precision, therefore we need to run the flow on the CPU
736         fp64_on_cpu = True
737         print(f"Using device: {device}, performing fp64 on CPU: {fp64_on_cpu}")
738         model.to(device)

```

```

739     # ===== Load and train =====
740     spectra, labels, spectra_length, n_labels, labelNames, scaler =
get_normalized_data(DATA_PATH) # Load normalized the data
741
742     spectra_tensor = torch.tensor(spectra, dtype=torch.float32).unsqueeze(1) # Add a channel
dimension for the CNN
743     labels_tensor = torch.tensor(labels, dtype=torch.float32)
744
745     print(f"Spectra tensor shape: {spectra_tensor.shape}") # Should be (batch_size, 1,
16384)
746
747
748     train_loader, val_loader, test_loader = get_datasets(spectra_tensor, labels_tensor) #
Split the data into train, val, and test sets
749
750
751     model, train_losses, val_losses = train_nf_model(model, train_loader, val_loader, device,
pdf_model, num_epochs=epochs)
752
753     test_loss = evaluate_test_loss(model, test_loader, device)
754     print(f"Test Loss: {test_loss:.6f}")
755     plot_loss(train_losses, val_losses, test_loss, pdf_model, filename="loss_plot.png") #
Plot the loss, including test loss
756
757
758     # =====Visualize the PDF for a batch of input data =====
759     model.eval() # Set the model to evaluation mode
760     with torch.no_grad():
761         for batch_inputs, batch_labels in test_loader:
762             batch_inputs = batch_inputs.to(device)
763             print(batch_inputs.shape)
764             model.visualize_pdf(
765                 input_data=batch_inputs,
766                 filename="pdf_visualization.png",
767                 pdf_model=pdf_model,
768                 samplesize=10000,
769                 batch_index=0,
770                 truth=batch_labels[0] if batch_labels is not None else None
771             )
772             break # Visualize only the first batch
773
774
775     compute_and_plot_coverage(model, test_loader, device, scaler, pdf_model)
776
777
778
779
780
781
782     # ===== End of Code =====
783

```

