

Exercise 5: Generative Adversarial Networks

Erik Stolt

What you did and how

Generative Adversarial Networks (GANs) work by incorporating two neural network architectures: a generator (G) and a discriminator (D). The generator learns to produce realistic images from an abstract input distribution, typically Gaussian noise. The discriminator learns to determine whether an image is real, meaning from the training dataset, or fake, meaning generated by G.

The discriminator's loss function is the binary cross-entropy (BCE) loss. Real images are labeled as 1 and fake images as 0. During training, the discriminator is presented with both real and fake images in each batch, along with their corresponding labels. It updates its weights to improve its ability to classify them correctly. The discriminator's total loss is computed as the average of the BCE loss on real and fake samples.

The generator is trained indirectly through the discriminator. To train the generator, fake images are generated and passed to the discriminator, but this time with the label 1, as if they were real. The generator's objective is to produce fake images that the discriminator classifies as real. When the discriminator is fooled by a fake image and predicts a value close to 1, the generator receives a lower loss. Therefore, the generator learns to generate fake images that closely resemble real images from the training data.

For this exercise, we defined a generator and a discriminator network, and trained them in an adversarial setup where the generator learned to produce increasingly realistic digit images while the discriminator tried to distinguish them from real MNIST digits. We monitored the training progress by generating sample images at different epochs by fixing the noise beforehand, which clearly showed how the image quality improved over time. In practice we first normalized the input images s.t. the pixel values transformed like $[0, 1] \rightarrow [-1, 1]$ which we also later denormalized for visualization according to $\text{image} = \text{image} * 0.5 + 0.5$. Regarding the network structures for the generator and discriminator, we settled on two standard NN. The generator had tanh activation in the output layer to transform the values to the $[-1, 1]$ range. The discriminator instead used sigmoid activation since its task is a binary classification task. The networks were then trained for 30 epochs with a batch size of 32, monitoring the performances using Tensorboard.

What results you obtained

Example print-out during training:

```
Epoch [28/30] | Loss D: 0.6806 | lr D: 5e-05 | Loss G: 0.7323 | lr G: 0.0001 | 135.27s
Epoch [29/30] | Loss D: 0.6682 | lr D: 5e-05 | Loss G: 0.7079 | lr G: 0.0001 | 122.10s
Epoch [30/30] | Loss D: 0.6654 | lr D: 5e-05 | Loss G: 0.7033 | lr G: 0.0001 | 126.16s
```

I also present the evolution of the digits and the loss as a function of epoch:




Figure 1: Loss vs epoch for both G and D .




Figure 2: Evolution of G generating digits.

We clearly see that the loss stays the same but the digits obviously get better. This is the result of a good trade-off between G and D improving at the same rate.

What challenges you encountered and what could be improved

GANs are notoriously unstable which lead to the discriminant having zero loss, hence, no learning. This is fixed by setting different initial lrs for the two models. Furthermore, I also replaced the true labels ones and zeros with 0.9 and 0.1 (instead of 1 and 0) in order to trick D not to be overconfident. This works because it

penalizes the model in being overconfident (see Appendix 1). I also added dropout and batch normalization in order for G not to be deterministic as this also lead to model collapse. Batch normalization is a critical part of G but often left out in D since it can cause numerical instabilities [1]. To improve the model one could use WGANs to get even better results and prevent model collapse.

This exercise also taught me a lot about interpreting losses. In GANs, monitoring losses is not enough and even misleading. If one only were to monitor the loss, model collapse seems good since it happens when the discriminators loss is zero. However, we also have the two scenarios

- Discriminator loss may go to zero when it is dominating, but that means G is not learning.
- Generator loss may go up even when it is improving image quality because D is getting better.

This means that in adversarial setups, both networks are playing a game, so, their losses do not converge in the usual sense.

Link to GitHub

Link: <https://github.com/erst6955/Advanced-Applied-Deep-Learning-in-Physics-And-Engineering>

Appendix 1

BCE is given by

$$\text{BCE}(p, y) = -(y \cdot \log(p) + (1 - y) \cdot \log(1 - p)),$$

which means that if $p = 0.99$ and $y = 1$ we would get $\text{BCE} \approx 0.01$ which is very small loss, hence, the gradient would vanish. However, if we let $y \rightarrow 0.9$ we would instead get $\text{BCE} \approx 0.47$ which is higher! Therefore, more learning is possible with this small change. The same goes for label $y \rightarrow 0.1$. This is a common trick which works well when D is learning too strong.

Code (for backup)

~\OneDrive\Desktop\Exercise 5 - Generative Adversarial Networks\GANs for MNIST.py

```

1
2
3 import torchvision
4 from torchvision import transforms, datasets
5 import torch
6 import torch.nn as nn
7 import torch.optim as optim
8 from torch.utils.data import DataLoader
9 from torch.utils.tensorboard import SummaryWriter
10 import matplotlib.pyplot as plt
11 import time
12 import random
13
14
15
16 # Hyperparameters
17 device = "cuda" if torch.cuda.is_available() else "cpu"
18 lr = 3e-4
19 batchSize = 32
20 logStep = 625
21 latent_dimension = 128
22 image_dimension = 784
23
24 myTransforms = transforms.Compose([ # we define a tranform that converts the image to
    tensor and normalizes it with mean and std of 0.5
25     transforms.ToTensor(),          # which will convert the image range from [0, 1] to
    [-1, 1]
26     transforms.Normalize((0.5,), (0.5,))
27 ])
28
29 dataset = datasets.MNIST(root="dataset/", transform=myTransforms, download=True)
30 loader = DataLoader(dataset, batch_size=batchSize, shuffle=True)
31
32
33 class Generator(nn.Module):
34     """
35     Generator Model which takes a random noise vector and generates a fake image. The
    input noise in latent space is an abstract space where each point corresponds to a
    different kind of image.
36     The noise vector doesn't directly mean anything in the beginning, but after training,
    it gets mapped by the generator to a specific kind of image – like a "7", or a "2", or a
    "3". We make
37     assumptions about how the input noise is distributed.
38     """
39     def __init__(self):
40         super().__init__()
41         self.gen = nn.Sequential( # simple NN
42             nn.Linear(latent_dimension, 256),
43             nn.BatchNorm1d(256),
44             nn.ReLU(),
45
46             nn.Linear(256, 512),

```

```

47         nn.BatchNorm1d(512),
48         nn.ReLU(),
49
50         nn.Linear(512, 1024),
51         nn.BatchNorm1d(1024),
52         nn.ReLU(),
53
54         nn.Linear(1024, image_dimension),
55         nn.Tanh(), # tanh activation function to get the output in the range of [-1,
1]
56     )
57
58     def forward(self, x):
59         return self.gen(x)
60
61 class Discriminator(nn.Module):
62     """
63     Discriminator Model which takes an image and outputs a probability of it being real or
64     fake. Furthermore,
65     the discriminator is a binary classifier which uses the sigmoid activation and the BCE
66     loss function.
67     """
68     def __init__(self):
69         super().__init__()
70         self.disc = nn.Sequential(
71             nn.Linear(image_dimension, 1024),
72             nn.LeakyReLU(0.2),
73             nn.Dropout(0.3), # add this
74             nn.Linear(1024, 512),
75             nn.LeakyReLU(0.2),
76             nn.Dropout(0.3), # add this
77             nn.Linear(512, 256),
78             nn.LeakyReLU(0.2),
79             nn.Linear(256, 1),
80             nn.Sigmoid(),
81         )
82
83     def forward(self, x):
84         return self.disc(x)
85
86 def train_models(discriminator, generator, fixed_noise, num_examples, opt_discriminator,
87 opt_generator, criterion, epochs, writer):
88     """
89     Trains the GAN models using the MNIST dataset.
90     The generator uses BCE where 1 is the label for fake images and 0 is the label for
91     real images. This mean that it learns
92     to generate fake images that are similar to the real images in the dataset since the
93     discriminator has real images labeled as
94     1 and fake images labeled as 0. I.e., the generator aims to minimize the loss by
95     producing images that are classified as real by the discriminator.
96     The discriminator on the other hand is simply trained to identify the real and fake
97     images.
98     """

```

```

94     step = 0 # for tensorboard logging
95     gen_losses = []
96     disc_losses = []
97
98     # ===== Training Loop =====
99     for epoch in range(epochs):
100
101         start_time = time.time()
102
103         for batch_idx, (real, _) in enumerate(loader):
104             real = real.view(-1, image_dimension).to(device) # flatten the image in order
to pass it to the discriminator
105             batch_size = real.shape[0] # get the batch size
106
107             noise = torch.randn(batch_size, latent_dimension).to(device) # generate random
noise from a normal distribution
108             fake = generator(noise) # generate fake images from the noise
109
110             # Discriminator Loss
111             disc_real = discriminator(real).view(-1) # get the discriminator output for
real images
112             loss_real = criterion(disc_real, torch.full_like(disc_real, 0.9)) # pass
through BCE where real images are labeled as 1
113
114             disc_fake = discriminator(fake.detach()).view(-1) # get the discriminator
output for fake images
115             loss_fake = criterion(disc_fake, torch.full_like(disc_fake, 0.1)) # pass
through BCE where fake images are labeled as 0
116
117             loss_discriminator = (loss_real + loss_fake) / 2 # average the loss for real
and fake images
118
119             discriminator.zero_grad() # zero the gradients
120             loss_discriminator.backward(retain_graph=True) # backpropagate the loss
121             opt_discriminator.step() # update the discriminator weights
122
123             # ===== Generator Loss =====
124             # The generator tries to fool the discriminator, so we want the discriminator
to think that the fake images are real
125             output = discriminator(fake).view(-1) # get the discriminator output for fake
images
126             loss_generator = criterion(output, torch.ones_like(output)) # pass through BCE
where fake images are labeled as 1
127             # The generator tries to maximize the probability of the discriminator being
wrong
128
129             generator.zero_grad()
130             loss_generator.backward()
131             opt_generator.step()
132
133             if batch_idx % logStep == 0: # tensorboard logging
134                 with torch.no_grad():
135                     fake_images = generator(fixed_noise).reshape(-1, 1, 28, 28) # reshape
the fake images to 1 channel and 28x28

```

```

136         real_images = real.reshape(-1, 1, 28, 28) # reshape the real images to
137         1 channel and 28x28
138         imgGridFake = torchvision.utils.make_grid(fake_images, normalize=True)
139         # make a grid of fake images
140         imgGridReal = torchvision.utils.make_grid(real_images, normalize=True)
141         # make a grid of real images
142
143         # Denormalize: [-1, 1] → [0, 1]
144         fake_images = denormalize(fake_images)
145         real_images = denormalize(fake_images)
146
147         writer.add_image("MNIST Fake Images", imgGridFake, global_step=step)
148         writer.add_image("MNIST Real Images", imgGridReal, global_step=step)
149         writer.add_scalar("Loss Discriminator", loss_discriminator.item(),
150         step)
151         writer.add_scalar("Loss Generator", loss_generator.item(), step)
152
153         step += 1
154         for param_group in opt_generator.param_groups:
155             lr_g = param_group['lr']
156         for param_group in opt_discriminator.param_groups:
157             lr_d = param_group['lr']
158
159         # At the end of each epoch
160         gen_losses.append(loss_generator.item())
161         disc_losses.append(loss_discriminator.item())
162         elapsed = time.time() - start_time
163         print(f"Epoch [{epoch+1}/{epochs}] | Loss D: {loss_discriminator:.4f} | lr D:
164         {lr_d} | Loss G: {loss_generator:.4f} | lr G: {lr_g} | {elapsed:.2f}s")
165
166         get_fake_images(generator, fixed_noise, latent_dimension, num_examples, epoch) #
167         get some fake images for the epoch
168
169         return gen_losses, disc_losses
170
171
172 def denormalize(images):
173     return images * 0.5 + 0.5
174
175
176 def get_fake_images(generator, fixed_noise, latent_dimension, num_examples, epoch):
177     """
178     Returns a batch of fake images from the generator.
179     """
180
181     generator.eval()
182
183     # Generate images from noise
184     with torch.no_grad():
185         generated_images = generator(fixed_noise).reshape(-1, 1, 28, 28)
186         generated_images = denormalize(generated_images) # denormalize the images
187

```

```

184     # Create a grid for visualization
185     grid = torchvision.utils.make_grid(generated_images.cpu(), nrow=4, normalize=True)
186
187     # Plot the generated images
188     plt.figure(figsize=(8, 8))
189     plt.title(f"Fake Images Generated by GAN at epoch {epoch+1}")
190     plt.axis("off")
191     plt.imshow(grid.permute(1, 2, 0).squeeze())
192     plt.savefig(f"Figures/fake_images_{epoch+1}.png", bbox_inches='tight')
193     plt.close()
194
195
196
197 def plot_losses(gen_losses, disc_losses):
198     plt.figure(figsize=(10, 5))
199     plt.plot(gen_losses, label="Generator Loss")
200     plt.plot(disc_losses, label="Discriminator Loss")
201     plt.xlabel("Batch")
202     plt.ylabel("Loss")
203     plt.title("Generator and Discriminator Loss During Training")
204     plt.legend()
205     plt.grid(True)
206     plt.tight_layout()
207     plt.savefig("Figures/loss_plot.png", dpi=300)
208
209
210
211 # ===== Run Code
212 =====
213 discriminator = Discriminator().to(device)
214 generator = Generator().to(device)
215
216 lr_G = 1e-4
217 lr_D = 5e-5
218
219 opt_generator = optim.Adam(generator.parameters(), lr=lr_G, betas=(0.5, 0.999))
220 opt_discriminator = optim.Adam(discriminator.parameters(), lr=lr_D, betas=(0.5, 0.999))
221 criterion = nn.BCELoss() # Binary Cross Entropy Loss
222
223 epochs = 30
224 num_examples = 3
225
226 fixed_noise = torch.randn(num_examples, latent_dimension).to(device) # fixed noise for
227 reproducibility after every epoch
228 writer = SummaryWriter("runs/GAN_MNIST")
229
230 gen_losses, disc_losses = train_models(discriminator, generator, fixed_noise,
231 num_examples, opt_discriminator, opt_generator, criterion, epochs, writer)
232 plot_losses(gen_losses, disc_losses)
233
234 writer.close()
235 print("Complete!")
236
237
238

```


235
236
237
238
239
240

References

- [1] Alec Radford, Luke Metz, and Soumith Chintala. “Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks”. In: *arXiv preprint arXiv:1511.06434* (2016). URL: <https://arxiv.org/abs/1511.06434>.