

Exercise 6: Simple Diffusion

Erik Stolt

What you did and how

In this exercise we explore Diffusion Model, a relatively simple architecture that exploits noise for data augmentation and to generate new samples following the same distribution as the training data. In its simplest form, we start by, in time steps, adding noise to a training data point (this is the forward pass), then we train a NN to denoise the the training example in order to reconstruct the image. In the end, the idea is that we can simply start with an image that is complete noise (Gaussian in our case, but it could in theory be another distribution), and then pass it through the network to get a new unique image that does not appear in the data set. In practice we do this by following the algorithm [1]:

Algorithm 1 Training

- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on

$$\nabla_{\theta} \left\| \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta} \left(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t \right) \right\|^2$$

- 6: **until** converged
-

where \mathbf{x}_0 is an image from the training data set. The advantage of diffusion models is that a single training example can be used multiple times, since in each iteration we add an independent noise vector to it. Although in this case, our image is only a single pixel. Notice that $\boldsymbol{\epsilon}_{\theta}$ is actually the total noise added to the image over t time steps. This is not what we want, we want the model to predict the noise added to the image in each time step. It turns out, as done in the paper, that one can parametrize the model using a change of variable to obtain this. The end result is given in algorithm 2 [1] where we now have an additional scaling factor in front of $\boldsymbol{\epsilon}_{\theta}$:

Algorithm 2 Sampling

- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 2: **for** $t = T, \dots, 1$ **do**
 - 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
 - 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
 - 5: **end for**
 - 6: **return** \mathbf{x}_0
-

We can now sample from algorithm 2 as man times as we want to end up with a distribution that follows the training data, and which we can input random noise \mathbf{x}_T to and receive a new image.

What results you obtained

One example print-out during training looked like this:

Epoch 77/1000 | Training loss: 0.35963864673693213 | Validation Loss: 1.7202495336532593

And the following plots were obtained

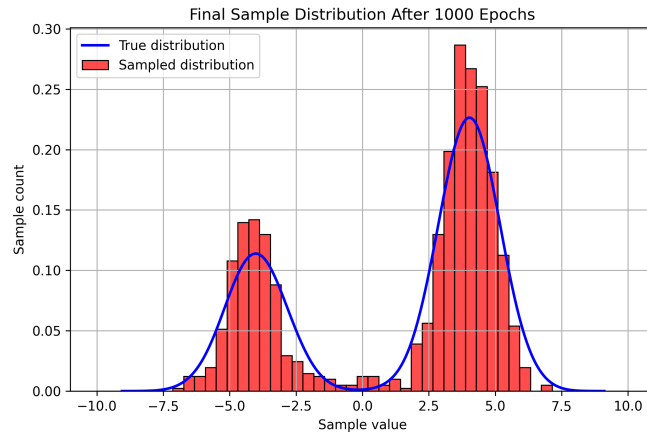


Figure 1: Posterior distribution of the predicted pixel values (red). Notice that it well aligned with the prior distribution of the training data (blue).

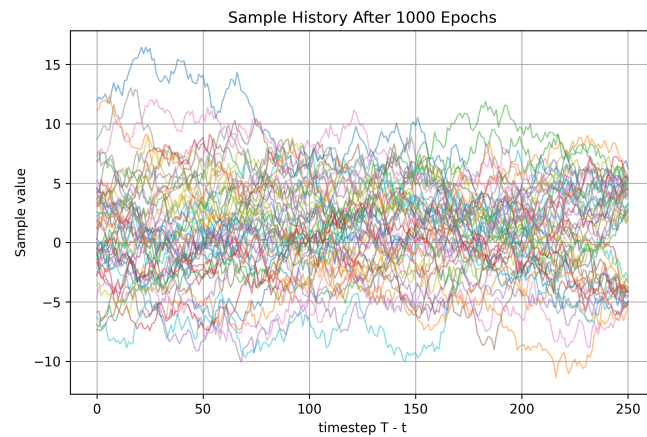


Figure 2: Evolution of the predicted pixel value as a function of time. Notice that at $t = 250$, the trajectories has converged to the pdf illustrated in Fig.1

Overall, the results are as expected.

What challenges you encountered and what could be improved

The main problem I encountered in this exercise was to correctly handle the shapes of all the vectors. I took me a while to correctly implement the algorithms in a way that the code could handle. As this is a relatively simple exercise, there is not much to be improved since a simple network captured all the core features. If that was not the case one would need a better NN architecture. One can always implement this in order to obtain even better results, although, this is not needed here.

Link to GitHub

Link: <https://github.com/erst6955/Advanced-Applied-Deep-Learning-in-Physics-And-Engineering>

Code (for backup)

~\OneDrive\Desktop\Exercise 6 - Simple Diffusion\Simple diffusion.py

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 import torch
4 import seaborn as sns # a useful plotting library on top of matplotlib
5 from tqdm.auto import tqdm # a nice progress bar
6
7
8 def normalize(x, mean, std):
9     return (x - mean) / std
10
11 def denormalize(x, mean, std):
12     return x * std + mean
13
14
15 # generate a dataset of 1D data from a mixture of two Gaussians
16 # this is a simple example, but you can use any distribution
17 data_distribution = torch.distributions.mixture_same_family.MixtureSameFamily(
18     torch.distributions.Categorical(torch.tensor([1, 2])),
19     torch.distributions.Normal(torch.tensor([-4., 4.]), torch.tensor([1., 1.]))
20 )
21
22 dataset = data_distribution.sample(torch.Size([10000])) # create training data set
23 dataset_validation = data_distribution.sample(torch.Size([1000])) # create validation data
24 set
25
26 mean = dataset.mean()
27 std = dataset.std()
28 dataset_norm = normalize(dataset, mean, std)
29 dataset_validation_norm = normalize(dataset_validation, mean, std)
30
31 # ===== HYPERPARAMETERS =====
32
33 TIME_STEPS = 250
34 BETA = torch.full((TIME_STEPS,), 0.02)
35 N_EPOCHS = 1000
36 BATCH_SIZE = 64
37 LEARNING_RATE = 0.8e-4
38
39 # define the neural network that predicts the amount of noise that was
40 # added to the data
41 # the network should have two inputs (the current data and the time step)
42 # and one output (the predicted noise)
43 # ===== Model
44 =====
45 class NoisePredictor(torch.nn.Module):
46     def __init__(self): # define simple nn with concatenation
47         super(NoisePredictor, self).__init__()
48         self.fc1 = torch.nn.Linear(2, 128) # Input layer (data + time step)
49         self.fc2 = torch.nn.Linear(128, 128) # Hidden layer
50         self.fc3 = torch.nn.Linear(128, 1) # Output layer (predicted noise)
51         self.tanh = torch.nn.Tanh()

```

```

51
52     def forward(self, x, t):
53         # Concatenate data and time step
54         t = t.unsqueeze(1) # [BATCH SIZE, 1]
55         x = x.view(x.size(0), -1) # Flatten the input data s.t. [BATCH SIZE, 1]
56
57         input_tensor = torch.cat((x, t.float()), dim=1) # [BATCH SIZE, 2] e.g. Sample 1
→ x_t = 0.34, timestep t = 5
58
59         x = self.tanh(self.fc1(input_tensor))
60         x = self.tanh(self.fc2(x))
61         x = self.fc3(x)
62         return x
63
64
65     def train_model(g, dataset_norm, dataset_validation_norm):
66
67         epochs = tqdm(range(N_EPOCHS)) # this makes a nice progress bar
68         criterion = torch.nn.MSELoss() # Use Mean Squared Error Loss
69         optimizer = torch.optim.Adam(g.parameters(), lr=LEARNING_RATE)
70
71         bar_alpha = torch.cumprod(1 - BETA, dim=0) # Precompute the cumulative product for all
time steps
72         total_loss = 0
73         n_batches = 0
74
75         for e in epochs: # loop over epochs
76             g.train()
77             # loop through batches of the dataset, reshuffling it each epoch
78             indices = torch.randperm(dataset_norm.shape[0]) # shuffle the dataset
79             shuffled_dataset_norm = dataset_norm[indices] # shuffle the dataset
80
81             for i in range(0, shuffled_dataset_norm.shape[0] - BATCH_SIZE, BATCH_SIZE): # loop
through the dataset in batches
82                 x0 = shuffled_dataset_norm[i:i + BATCH_SIZE].view(-1, 1) # sample a batch of
data and add dimension [B] --> [B,1] since this is necessary format for the NN
83
84                 # here, implement algorithm 1 of the DDPM paper
(https://arxiv.org/abs/2006.11239)
85                 t = torch.randint(0, TIME_STEPS, (BATCH_SIZE,)) # sample uniformly a time
step
86                 noise = torch.randn_like(x0) # sample the noise
87                 bar_alpha_t = bar_alpha[t].view(-1, 1) # compute the product of alphas up to
time t and add dimension
88
89                 x_t = torch.sqrt(bar_alpha_t) * x0 + torch.sqrt(1 - bar_alpha_t) * noise # -->
[B, 1]
90                 predicted_noise = g(x_t, t.float()) # compute the predicted noise
91
92                 # compute the loss (mean squared error between predicted noise and true noise)
93                 loss = criterion(predicted_noise, noise)
94
95                 # backpropagation and loss stuff
96                 optimizer.zero_grad()
97                 loss.backward()

```

```

98         optimizer.step()
99
100         total_loss += loss.item()
101         n_batches += 1
102         avg_loss = total_loss / n_batches
103
104         # compute the loss on the validation set
105         g.eval()
106         with torch.no_grad():
107             x0 = dataset_validation_norm
108             t = torch.randint(0, TIME_STEPS, (x0.shape[0],)) # sample a time step for
validation
109             noise = torch.randn_like(x0) # sample the noise
110             val_bar_alpha_t = bar_alpha[t] # compute the product of alphas up to time t
111             x_t = torch.sqrt(val_bar_alpha_t) * x0 + torch.sqrt(1 - val_bar_alpha_t) *
noise # add noise to the validation data
112
113             predicted_noise = g(x_t, t.float())# Compute the predicted noise
114
115             val_loss = criterion(predicted_noise, noise) # Calculate the validation loss
116             print(f" Epoch {e+1}/{N_EPOCHS}| Training loss: {avg_loss} | Validation Loss:
{val_loss.item()}")
117
118
119
120 def sample_and_track(g, count):
121     """
122     Sample from the model by applying the reverse diffusion process
123
124     Here, implement algorithm 2 of the DDPM paper (https://arxiv.org/abs/2006.11239)
125
126     Parameters
127     -----
128     g : torch.nn.Module
129         The neural network that predicts the noise added to the data
130     count : int
131         The number of samples to generate in parallel
132
133     Returns
134     -----
135     x : torch.Tensor
136         The final sample from the model
137     -----
138     Perform reverse diffusion:
139     - Return final sampled values for all `count`
140     - Track one sample (sample 0) over time using x_batch
141     """
142     g.eval()
143     bar_alpha = torch.cumprod(1 - BETA, dim=0)
144
145     x_batch = torch.randn(count, 1) # [count, 1]
146     tracked_index = 0 # track first index
147     history = [x_batch[tracked_index].item()] # Track first sample
148

```

```

149     for t in range(TIME_STEPS - 1, -1, -1):
150         t_tensor_batch = torch.full((count,), t, dtype=torch.long)
151
152         # Predict noise
153         pred_noise_batch = g(x_batch, t_tensor_batch.float()).view(-1, 1)
154
155         # Get scalars
156         bar_alpha_t = bar_alpha[t]
157         alpha_t = 1 - BETA[t]
158         factor = (1 - alpha_t) / torch.sqrt(1 - bar_alpha_t)
159         sigma_t = torch.sqrt(BETA[t])
160
161         # Random noise (zero for last step)
162         z_batch = torch.randn_like(x_batch) if t > 0 else torch.zeros_like(x_batch)
163
164         # Reverse step
165         x_batch = (1 / torch.sqrt(alpha_t)) * (x_batch - factor * pred_noise_batch) +
sigma_t * z_batch # Posterior decoded pixel value
166
167         # Track sample 0
168         history.append(x_batch[tracked_index].item())
169
170     # Denormalize
171     samples = denormalize(x_batch, mean, std).detach().numpy().flatten()
172     history = denormalize(torch.tensor(history), mean, std).numpy()
173
174     return samples, history
175
176 # ===== Plots =====
177 def plot_distribution(samples):
178     fig, ax = plt.subplots(1, 1, figsize=(8, 5))
179     bins = np.linspace(-10, 10, 50)
180     sns.kdeplot(dataset.numpy().flatten(), ax=ax, color='blue', label='True distribution',
linewidth=2)
181     sns.histplot(samples, ax=ax, bins=bins, color='red', label='Sampled distribution',
stat='density', alpha=0.7)
182     ax.legend()
183     ax.set_xlabel('Sample value')
184     ax.set_ylabel('Sample count')
185     plt.title(f"Final Sample Distribution After {N_EPOCHS} Epochs")
186     plt.grid(True)
187
188     plt.savefig("Figures/final_distribution.png", dpi=300)
189     plt.close()
190
191
192 def plot_monte_carlo(all_histories):
193     plt.figure(figsize=(8, 5))
194     for history in all_histories:
195         plt.plot(range(TIME_STEPS + 1), history, alpha=0.5, linewidth=1)
196     plt.xlabel('timestep T - t')
197     plt.ylabel('Sample value')
198     plt.title(f'Sample History After {N_EPOCHS} Epochs')
199     plt.grid(True)

```

```
200
201     plt.savefig("Figures/sample_history.png", dpi=300)
202     plt.close()
203
204
205 def generate_plots(g, N):
206     all_histories = []
207
208     for i in range(N):
209         samples, history = sample_and_track(g, 1000)
210         # Only save histogram plot for first run
211         if i == 0:
212             plot_distribution(samples)
213
214         all_histories.append(history)
215
216     plot_monte_carlo(all_histories)
217
218
219 # ===== RUNNING THE CODE =====
220
221 g = NoisePredictor()
222 train_model(g, dataset_norm, dataset_validation_norm)
223
224 generate_plots(g, 50)
225
226
227
228
229
230
231
232
233
234
```

References

- [1] Tom B Brown et al. “Language Models are Few-Shot Learners”. In: *arXiv preprint arXiv:2006.11239* (2020). URL: <https://arxiv.org/abs/2006.11239>.