

~\OneDrive\Desktop\Exercise 6 - Simple Diffusion\Simple diffusion.py

```

1  import matplotlib.pyplot as plt
2  import numpy as np
3  import torch
4  import seaborn as sns # a useful plotting library on top of matplotlib
5  from tqdm.auto import tqdm # a nice progress bar
6
7
8  def normalize(x, mean, std):
9      return (x - mean) / std
10
11 def denormalize(x, mean, std):
12     return x * std + mean
13
14
15 # generate a dataset of 1D data from a mixture of two Gaussians
16 # this is a simple example, but you can use any distribution
17 data_distribution = torch.distributions.mixture_same_family.MixtureSameFamily(
18     torch.distributions.Categorical(torch.tensor([1, 2])),
19     torch.distributions.Normal(torch.tensor([-4., 4.]), torch.tensor([1., 1.])))
20 )
21
22 dataset = data_distribution.sample(torch.Size([10000])) # create training data set
23 dataset_validation = data_distribution.sample(torch.Size([1000])) # create validation data
24 set
25
26 mean = dataset.mean()
27 std = dataset.std()
28 dataset_norm = normalize(dataset, mean, std)
29 dataset_validation_norm = normalize(dataset_validation, mean, std)
30
31 # ===== HYPERPARAMETERS =====
32
33 TIME_STEPS = 250
34 BETA = torch.full((TIME_STEPS,)), 0.02)
35 N_EPOCHS = 1000
36 BATCH_SIZE = 64
37 LEARNING_RATE = 0.8e-4
38
39 # define the neural network that predicts the amount of noise that was
40 # added to the data
41 # the network should have two inputs (the current data and the time step)
42 # and one output (the predicted noise)
43 # ===== Model
44 =====
45 class NoisePredictor(torch.nn.Module):
46     def __init__(self): # define simple nn with concatenation
47         super(NoisePredictor, self).__init__()
48         self.fc1 = torch.nn.Linear(2, 128) # Input layer (data + time step)
49         self.fc2 = torch.nn.Linear(128, 128) # Hidden layer
50         self.fc3 = torch.nn.Linear(128, 1) # Output layer (predicted noise)
51         self.tanh = torch.nn.Tanh()

```

```

51
52     def forward(self, x, t):
53         # Concatenate data and time step
54         t = t.unsqueeze(1) # [BATCH SIZE, 1]
55         x = x.view(x.size(0), -1) # Flatten the input data s.t. [BATCH SIZE, 1]
56
57         input_tensor = torch.cat((x, t.float()), dim=1) # [BATCH SIZE, 2] e.g. Sample 1
→ x_t = 0.34, timestep t = 5
58
59         x = self.tanh(self.fc1(input_tensor))
60         x = self.tanh(self.fc2(x))
61         x = self.fc3(x)
62         return x
63
64
65     def train_model(g, dataset_norm, dataset_validation_norm):
66
67         epochs = tqdm(range(N_EPOCHS)) # this makes a nice progress bar
68         criterion = torch.nn.MSELoss() # Use Mean Squared Error Loss
69         optimizer = torch.optim.Adam(g.parameters(), lr=LEARNING_RATE)
70
71         bar_alpha = torch.cumprod(1 - BETA, dim=0) # Precompute the cumulative product for all
time steps
72         total_loss = 0
73         n_batches = 0
74
75         for e in epochs: # loop over epochs
76             g.train()
77             # loop through batches of the dataset, reshuffling it each epoch
78             indices = torch.randperm(dataset_norm.shape[0]) # shuffle the dataset
79             shuffled_dataset_norm = dataset_norm[indices] # shuffle the dataset
80
81             for i in range(0, shuffled_dataset_norm.shape[0] - BATCH_SIZE, BATCH_SIZE): # loop
through the dataset in batches
82                 x0 = shuffled_dataset_norm[i:i + BATCH_SIZE].view(-1, 1) # sample a batch of
data and add dimension [B] --> [B,1] since this is necessary format for the NN
83
84                 # here, implement algorithm 1 of the DDPM paper
(https://arxiv.org/abs/2006.11239)
85                 t = torch.randint(0, TIME_STEPS, (BATCH_SIZE,)) # sample uniformly a time
step
86                 noise = torch.randn_like(x0) # sample the noise
87                 bar_alpha_t = bar_alpha[t].view(-1, 1) # compute the product of alphas up to
time t and add dimension
88
89                 x_t = torch.sqrt(bar_alpha_t) * x0 + torch.sqrt(1 - bar_alpha_t) * noise # -->
[B, 1]
90                 predicted_noise = g(x_t, t.float()) # compute the predicted noise
91
92                 # compute the loss (mean squared error between predicted noise and true noise)
93                 loss = criterion(predicted_noise, noise)
94
95                 # backpropagation and loss stuff
96                 optimizer.zero_grad()
97                 loss.backward()

```

```

98         optimizer.step()
99
100         total_loss += loss.item()
101         n_batches += 1
102         avg_loss = total_loss / n_batches
103
104         # compute the loss on the validation set
105         g.eval()
106         with torch.no_grad():
107             x0 = dataset_validation_norm
108             t = torch.randint(0, TIME_STEPS, (x0.shape[0],)) # sample a time step for
validation
109             noise = torch.randn_like(x0) # sample the noise
110             val_bar_alpha_t = bar_alpha[t] # compute the product of alphas up to time t
111             x_t = torch.sqrt(val_bar_alpha_t) * x0 + torch.sqrt(1 - val_bar_alpha_t) *
noise # add noise to the validation data
112
113             predicted_noise = g(x_t, t.float())# Compute the predicted noise
114
115             val_loss = criterion(predicted_noise, noise) # Calculate the validation loss
116             print(f" Epoch {e+1}/{N_EPOCHS}| Training loss: {avg_loss} | Validation Loss:
{val_loss.item()}")
117
118
119
120 def sample_and_track(g, count):
121     """
122     Sample from the model by applying the reverse diffusion process
123
124     Here, implement algorithm 2 of the DDPM paper (https://arxiv.org/abs/2006.11239)
125
126     Parameters
127     -----
128     g : torch.nn.Module
129         The neural network that predicts the noise added to the data
130     count : int
131         The number of samples to generate in parallel
132
133     Returns
134     -----
135     x : torch.Tensor
136         The final sample from the model
137     -----
138     Perform reverse diffusion:
139     - Return final sampled values for all `count`
140     - Track one sample (sample 0) over time using x_batch
141     """
142     g.eval()
143     bar_alpha = torch.cumprod(1 - BETA, dim=0)
144
145     x_batch = torch.randn(count, 1) # [count, 1]
146     tracked_index = 0 # track first index
147     history = [x_batch[tracked_index].item()] # Track first sample
148

```

```

149     for t in range(TIME_STEPS - 1, -1, -1):
150         t_tensor_batch = torch.full((count,), t, dtype=torch.long)
151
152         # Predict noise
153         pred_noise_batch = g(x_batch, t_tensor_batch.float()).view(-1, 1)
154
155         # Get scalars
156         bar_alpha_t = bar_alpha[t]
157         alpha_t = 1 - BETA[t]
158         factor = (1 - alpha_t) / torch.sqrt(1 - bar_alpha_t)
159         sigma_t = torch.sqrt(BETA[t])
160
161         # Random noise (zero for last step)
162         z_batch = torch.randn_like(x_batch) if t > 0 else torch.zeros_like(x_batch)
163
164         # Reverse step
165         x_batch = (1 / torch.sqrt(alpha_t)) * (x_batch - factor * pred_noise_batch) +
sigma_t * z_batch # Posterior decoded pixel value
166
167         # Track sample 0
168         history.append(x_batch[tracked_index].item())
169
170     # Denormalize
171     samples = denormalize(x_batch, mean, std).detach().numpy().flatten()
172     history = denormalize(torch.tensor(history), mean, std).numpy()
173
174     return samples, history
175
176 # ===== Plots =====
177 def plot_distribution(samples):
178     fig, ax = plt.subplots(1, 1, figsize=(8, 5))
179     bins = np.linspace(-10, 10, 50)
180     sns.kdeplot(dataset.numpy().flatten(), ax=ax, color='blue', label='True distribution',
linewidth=2)
181     sns.histplot(samples, ax=ax, bins=bins, color='red', label='Sampled distribution',
stat='density', alpha=0.7)
182     ax.legend()
183     ax.set_xlabel('Sample value')
184     ax.set_ylabel('Sample count')
185     plt.title(f"Final Sample Distribution After {N_EPOCHS} Epochs")
186     plt.grid(True)
187
188     plt.savefig("Figures/final_distribution.png", dpi=300)
189     plt.close()
190
191
192 def plot_monte_carlo(all_histories):
193     plt.figure(figsize=(8, 5))
194     for history in all_histories:
195         plt.plot(range(TIME_STEPS + 1), history, alpha=0.5, linewidth=1)
196     plt.xlabel('timestep T - t')
197     plt.ylabel('Sample value')
198     plt.title(f'Sample History After {N_EPOCHS} Epochs')
199     plt.grid(True)

```

```
200
201     plt.savefig("Figures/sample_history.png", dpi=300)
202     plt.close()
203
204
205 def generate_plots(g, N):
206     all_histories = []
207
208     for i in range(N):
209         samples, history = sample_and_track(g, 1000)
210         # Only save histogram plot for first run
211         if i == 0:
212             plot_distribution(samples)
213
214         all_histories.append(history)
215
216     plot_monte_carlo(all_histories)
217
218
219 # ===== RUNNING THE CODE =====
220
221 g = NoisePredictor()
222 train_model(g, dataset_norm, dataset_validation_norm)
223
224 generate_plots(g, 50)
225
226
227
228
229
230
231
232
233
234
```